

Building Blocks of a Microkernel: Lessons from seL4

Harshit Raj*

Department of Computer Science and Engineering
Indian Institute of Technology Kanpur

Abstract

In this report, we provide an analysis of the seL4 microkernel, exploring its building blocks and implementation. The seL4 microkernel is designed to minimize the Trusted Computing Base (TCB) and provide a secure system-building base by moving most kernel functionality to non-privileged user mode. This design enables the microkernel to be formally verified, providing an extra level of assurance in its security and reliability. We examine the seL4 kernel codebase and provide an overview of the kernel objects, components, and capabilities. Additionally, we present a design for building a very minimal microkernel. Overall, this report provides a detailed understanding of a microkernel and its building blocks.

Introduction

Kernel is a piece of code that serves as the core of an operating system. It manages the system's resources and provides a layer of abstraction between the hardware and software. The kernel is responsible for handling tasks such as memory management, process scheduling, device drivers, and input/output (I/O) operations. All these operations take place in a privileged mode, also known as kernel mode. On the other hand, programs operate in user mode which doesn't allow them to directly access hardware resources. Instead, they need to utilize the operating system to access these resources via different abstractions and interfaces.

Over the years this privileged code has increased with addition of newer technologies and this is as dangerous as useful. If any malicious access is allowed here that could compromise the entire system. This has occurred on many mainstream systems. Linux kernel comprises of 20 Million lines of source code, it is estimated that it contains tens of thousands of bugs [1]. This idea is captured by saying that Linux has a large trusted computing base (TCB), which is defined as the subset of the overall system that must be trusted to operate correctly for the system to be secure [2].

The idea behind a microkernel design is to drastically reduce the TCB to have a secure base to build a system. In a well-designed microkernel, such as seL4, it is of the order of ten thousand lines of source code. As seen in Figure 1 the monolithic kernel structure provides every functionality from kernel mode, i.e. privileged mode making the TCB large. On the other hand, the microkernel structure provides only the bare minimum functionality from kernel mode and the rest of the functionality is provided by user mode programs. This makes the TCB small and secure. However the downside of this approach is that it requires a lot of user mode programs to be written to provide the functionality and this incurs a performance penalty.

In this report we will be looking at the building blocks of a microkernel and how they are implemented in seL4. We will also be looking at the design to build a teaching microkernel using gem5 simulator [3]. This report aims to explore the concept of microkernels and their significance in designing secure operating systems. A microkernel design addresses this problem by minimizing the TCB to only include the essential functions of the kernel.

*Supervised by: Prof. Debadatta Mishra

Submitted: Apr 25, 2023

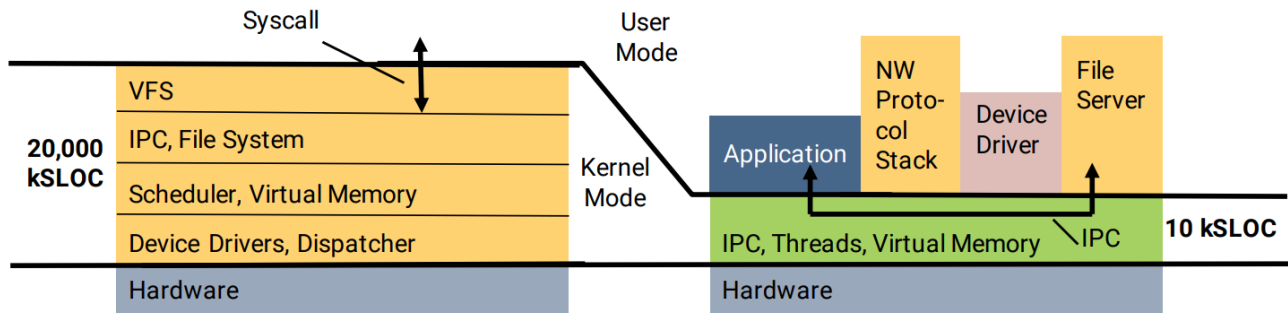


Figure 1: To left is a monolithic structure and microkernel structure on right [2]

Microkernel Design

The primary objective of Microkernels is to offer a small and secure kernel that is critical for security. The kernel provides only the essential functionality in kernel mode, and the remaining functionality is provided by user mode programs. This approach helps keep the Trusted Computing Base (TCB) small and secure, which reduces the attack surface. A well-designed microkernel, such as seL4, is typically only about ten thousand lines of source code, in contrast to the Linux kernel, which is approximately 20 million lines of code. This significant difference in size results in a much smaller TCB in a microkernel design, as shown in Figure 1. Let us now consider a microkernel design principle.

Minimality & Generality

The principles of minimality and Inter-Process Communication (IPC) performance were the primary drivers of Liedtke's designs. He firmly believed that minimality enhances IPC performance. To express this idea, he introduced the microkernel *minimality principle* as follows: [4]

A concept is tolerated inside the μ -kernel only if moving it outside the kernel, i.e. permitting competing implementations, would prevent the implementation of the system's required functionality

L4 had an implicit driver for the design of kernel mechanisms, which was generality: the aim was to create a foundation on which various systems could be built, almost anything that

could run on a processor powerful enough to provide protection.

However, none of the designers of L4 kernels to date claim that they have developed a "pure" microkernel that strictly adheres to the minimality principle. For instance, all of them have a scheduler in the kernel, which implements a specific scheduling policy, usually a hard-priority round-robin. To date, no one has been able to come up with a general in-kernel scheduler or a viable mechanism that delegates all scheduling policy to the user level without imposing a significant overhead.

In conclusion, microkernels maintain minimality as a crucial design principle and generality as the overall goal. [5]

Building Blocks

In this section, we will briefly examine the various building blocks that make up a microkernel.

Components Unlike traditional operating systems, microkernels do not provide any typical OS services. Instead, microkernel designs use **user-mode** programs called components to provide necessary system services. For example, a file system component provides file system services. These components are loaded into the kernel and run in user mode. The kernel provides the Inter-Process Communication (IPC) mechanism for these components to communicate with each other. Even the **user application** is a component in this design, as depicted in Figure 1.

Kernel Objects Microkernels also include a limited set of well-defined core kernel objects. These objects run in privileged mode and are respon-

sible for providing basic services, including an entry point for interrupt handling, IPC, basic virtual memory management, thread management, and scheduling. These are the only functionalities provided by the kernel. The rest of the services are provided by the components, as shown in Figure 1.

Bootng System As the components and kernel objects are complex and involved, they need to be configured before the system boots up and sent to appropriate locations in memory. This is the responsibility of the build system, which configures the components and kernel objects and sets metadata for the components.

We will discuss each of these building blocks in more detail later.

Components

Components are user-mode programs that provide system services and also function as application programs. The kernel provides an Inter-Process Communication (IPC) mechanism for these components to communicate with each other, if permitted.

Nature of Components:

- Components are user-mode programs and do not possess special privileges. All logic is implemented in user mode, and components do not have access to kernel mode.
- Components are independent of each other. If one component fails or is compromised, it does not affect the other components. The only impact is on the service provided by the failed component.
- Components communicate with each other using the IPC mechanism provided by the kernel.
- Components must be granted proper access rights at build time to access kernel objects, such as invoking IPC or accessing memory.
- Components are configurable, allowing them to provide different services, interrupt handling, etc., at build time.
- Components are reusable, meaning they can be used in different systems.

Representation of Components

Components are represented as executables using Thread Control Blocks (TCBs). TCBs are data structures that contain information about the component and are used by the kernel to manage the component's execution. TCBs store information needed for context switching, scheduling, priority, IPC buffer, IPC call, capabilities, etc.

Data in TCB The TCB contains various information, including:

- Architecture-specific TCB state, such as the program counter and stack pointer.
- Thread state: a three-word data structure indicating the current state of the component, which can be Inactive, Running, Restart, Blocked on Receive, Blocked on Send, Blocked on Reply, Blocked on Notification, or Idle Thread State.
- Current fault of the component, i.e., the fault that caused the component to be blocked.
- Maximum controlled priority of the component.
- Scheduling context of the component, which contains scheduling parameters such as core and timeslice.
- Scheduling and notification queues.

Capabilities

Capabilities are access rights granted to components to access kernel objects. These access rights are determined at build time and are represented as a data structure in the kernel. Capabilities are configurable at build time and are used to invoke IPC, access memory, etc. Capabilities are essential for implementing the principle of least privilege and determine which component can access which interrupt.

Kernel Objects

Kernel objects are fundamental components of the microkernel that provide essential services such as interrupt handling, inter-process communication (IPC), basic virtual memory management, thread management, and scheduling. These are the only functionalities that the kernel

provides directly. All other services are provided by components that incur an IPC overhead. The kernel objects are implemented in privileged mode. We have already seen one kernel object, the TCB, which is used to manage the execution of components. In this section, we will discuss the other kernel objects, namely system calls, IPC, virtual memory management, and interrupt handling.

System Calls

System calls are the means by which a user-level application interacts with the kernel. In a monolithic kernel, system calls are implemented as a set of functions that are called by the application. Each system call serves a specific purpose, such as opening a file, reading from a file, etc. In a microkernel, all these services are provided by components, and thus, only a limited number of system calls are needed.

Syscalls in seL4 and are needed in a microkernel are as follows:

- **Call:** This system call is used to invoke a function in a component. The component is identified by its TCB. The function to be invoked is identified by its index in the component's function table. The arguments to the function are passed in the IPC buffer.
- **ReplyRecv:** This system call is used to send a message to a component and wait for a reply. The component is identified by its TCB. The message is passed in the IPC buffer. The reply is also passed in the IPC buffer.
- **Send:** This system call is used to send a message to a component. The component is identified by its TCB. The message is passed in the IPC buffer.
- **NBSend:** This system call is used to send a message to a component. The component is identified by its TCB. The message is passed in the IPC buffer. This system call does not block the caller.
- **Recv:** This system call is used to receive a message from a component. The component is identified by its TCB. The message is passed in the IPC buffer.
- **NBRecv:** This system call is used to receive a message from a component. The compo-

nent is identified by its TCB. The message is passed in the IPC buffer. This system call does not block the caller.

- **Reply:** This system call is used to send a reply to a component. The component is identified by its TCB. The reply is passed in the IPC buffer.
- **Yield:** This system call is used to yield the CPU to another component.
- A microkernel may also need some thread management system calls, such as create thread, delete thread, etc.
- Syscalls are also implemented to manage virtual memory, such as map, unmap, etc. But a very elementary virtual memory management is implemented in the kernel. The rest of the virtual memory management is implemented in the components, accessed by a IPC call.

Inter Process Communication

Interprocess communication is possibly the most crucial and most discussed topic in microkernel as practically everything is an IPC. IPC is implemented as a Syscall in seL4. Components must be granted proper access rights at build time to invoke IPC to proper component. The IPC mechanism in seL4 is based on the following: IPC buffer, IPC endpoint, and IPC call.

IPC Buffer When a thread wants to send a message to another thread via IPC, it first needs to prepare a message in a buffer. The buffer is allocated from the thread's own memory space, and the message is copied into it. The sender then makes a system call to the kernel to initiate the IPC. The size of the buffer is typically small, which limits the size of messages that can be sent via IPC. However, this design choice makes the IPC mechanism more efficient, as there is no need for dynamic allocation of memory. The Kernel provides support for secure IPC communication by providing capabilities that allow only authorized threads to access a specific buffer.

Alternatively, the messages can be stored in a global buffer shared by all threads. This is the approach taken by the L4 microkernel. In this case, the kernel must ensure that the messages are not overwritten by other threads. This ap-

proach is more efficient than the previous one. However, it is less secure, as any thread can access the buffer.

IPC Endpoint IPC endpoint is a kernel object that represents a communication channel between two threads. It is the basic primitive used for inter-process communication (IPC) in seL4. The endpoint is used to send and receive messages between the threads. When a thread creates an endpoint, it gets a capability that allows it to access the endpoint. This capability can be used to send and receive messages through the endpoint. When a message is sent, the kernel copies the message into the endpoint's buffer and delivers it to the destination thread. When a message is received, the kernel copies the message from the endpoint's buffer into the receiver's buffer.

IPC Call IPC call is implemented as a system call. When a thread wants to initiate an IPC communication with another thread, it makes an IPC system call to the seL4 kernel. The IPC system call takes several arguments, including the capability to the endpoint, the message to be sent, and the length of the message. The kernel provides several different IPC system calls with different semantics, such as blocking and non-blocking variants, as well as variants that allow the calling thread to wait for multiple endpoints at once.

Virtual Memory Management

In the seL4 microkernel, virtual memory management is not provided beyond kernel primitives for manipulating hardware paging structures. Users must provide services for creating intermediate paging structures, as well as for mapping and unmapping pages.

Although users are free to define their own address space layout, there is one restriction: the seL4 kernel claims the high part of the virtual memory range, typically 0xe0000000 and above on most 32-bit platforms.

During the boot process, the seL4 kernel initializes the root task with a top-level hardware virtual memory object, referred to as a VSpace. Once all of the intermediate paging structures

have been mapped for a specific virtual address range, physical frames can be mapped into that range by invoking the frame capability.

Overall, the seL4 microkernel provides basic primitives for hardware-level virtual memory management, while leaving higher-level management to user-level services. This allows users to define their own virtual memory layout and implement customized management schemes, while retaining the security and performance benefits of the seL4 kernel.

Interrupts

In the seL4 microkernel, the root task is initially provided with a single capability that grants access to all interrupt (irq) numbers in the system. Components can register themselves to receive interrupts by communicating with the seL4 kernel, which manages the interrupt handling mechanism. When an interrupt occurs, the kernel sends an IPC message to the registered component. The component can then handle the interrupt and send a reply back to the kernel.

Boot System

Due to the complexity and interdependence of the components and kernel objects in the seL4 microkernel, they must be properly configured and placed in memory before the system can boot up. This task falls under the responsibility of the boot system, which is responsible for configuring the components and kernel objects, as well as setting metadata for the components.

The boot system ensures that all necessary components and kernel objects are properly linked and configured, and that the resulting binary image is compatible with the underlying hardware. In the microkernel, endpoints, capabilities, and interrupt handlers must be set up for all components while booting up. These components are responsible for handling various system services, and their proper initialization is critical for system stability and security. During boot seL4 also creates root server and idle server.

Root Server and Idle Server

The root server in seL4 microkernel provides a centralized resource management service for the entire system. It is responsible for initializing and managing the system's resources, including hardware devices, memory, and other system services. In addition to resource management, the root server is also responsible for initializing the system, and default interrupt handling.

The idle server, just like monolithic kernels, is a specialized thread that is responsible for managing system health and reducing power consumption. When no user-level process is scheduled for execution, the idle server remains idle, allowing the system to conserve power and reduce resource consumption.

Essential Blocks

View of a microkernel System

References

- [1] Simon Biggs, Damon Lee, and Gernot Heiser. "The Jury Is In: Monolithic OS Design Is Flawed". In: *Asia-Pacific Workshop on Systems (APSys)*. Korea: ACM SIGOPS, Aug. 2018. URL: https://ts.data61.csiro.au/publications/csiro_full_text//Biggs_LH_18.pdf.
- [2] Gernot Heiser. *The SeL4 Microkernel: An Introduction*. Tech. rep. 1.2. June 2020. URL: <https://sel4.systems/About/seL4-whitepaper.pdf>.
- [3] GEM5. *Gem5 Simulator*. <https://www.gem5.org/>. 2019.
- [4] Jochen Liedtke. "On μ -kernel Construction". In: *Proceedings of the 15th ACM Symposium on Operating Systems Principles*. Copper Mountain, CO, US: ACM, Oct. 1995, pp. 237–250. doi: 10.1145/224057.224075.
- [5] Gernot Heiser and Kevin Elphinstone. *L4 Microkernels: The Lessons from 20 Years of Research and Deployment*. NICTA and UNSW, Sydney, Australia. 2016. URL: https://trustworthy.systems/publications/nicta_full_text/8988.pdf.

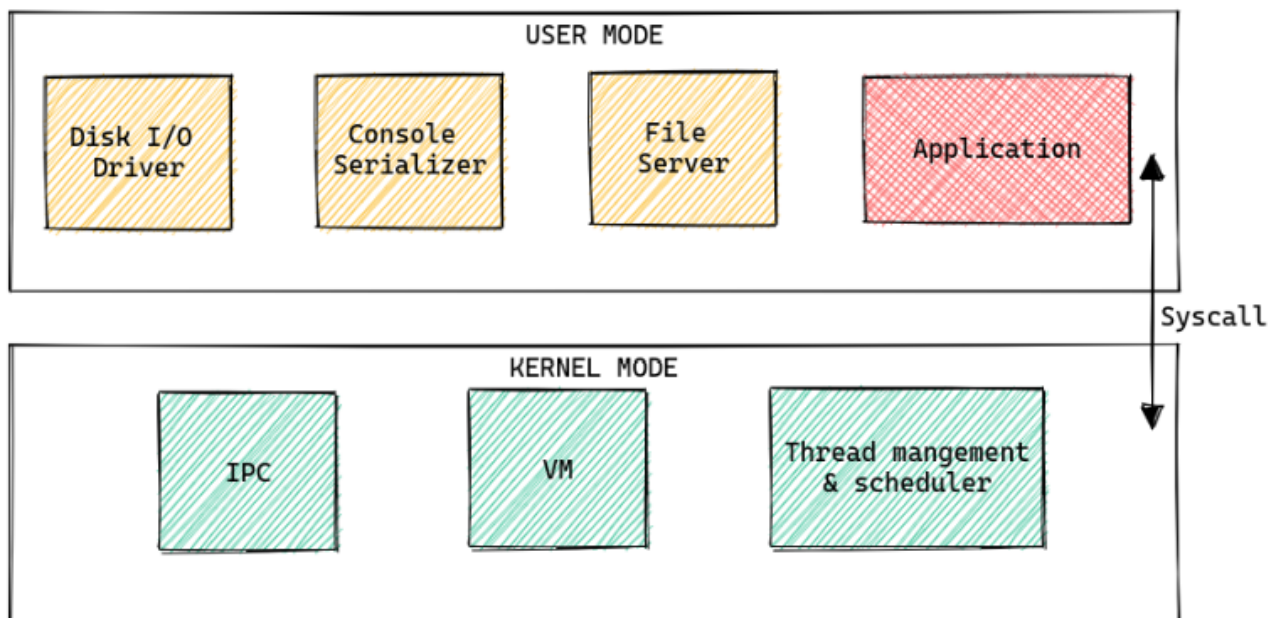


Figure 2: Example of an microkernel based system